

```
import datetime
import math
import os

import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow_datasets as tfds
import tensorflow_hub as hub

BATCH_SIZE = 32

train_data, train_info = tfds.load('cifar10', split='train[10%:90%]', with_info=True)
val_data = tfds.load('cifar10', split='train[0%:10%]')
test_data = tfds.load('cifar10', split='test')
print(train_data)

num_train_data = 0
for _ in train_data:
    num_train_data += 1
print(num_train_data)

num_val_data = 0
for _ in val_data:
    num_val_data += 1
print(num_val_data)

train_steps_per_epoch = math.ceil(num_train_data / BATCH_SIZE)
val_steps_per_epoch = math.ceil(num_val_data / BATCH_SIZE)

def normalizer(features, input_shape=[299, 299, 3], augment=True, seed=42):
    input_shape = tf.convert_to_tensor(input_shape)
    image = features['image']
    image = tf.image.convert_image_dtype(image, tf.float32)
    if augment:
        # Randomly applied horizontal flip:
        image = tf.image.random_flip_left_right(image, seed=seed)

        # Random B/S changes:
        image = tf.image.random_brightness(image, max_delta=0.1, seed=seed)
        image = tf.image.random_saturation(image, lower=0.5, upper=1.5, seed=seed)
        image = tf.clip_by_value(image, 0.0, 1.0) # keeping pixel values in check

    # Random resize and random crop back to expected size:

    random_scale_factor = tf.random.uniform([1], minval=1., maxval=1.4, dtype=tf.float32,
    scaled_height = tf.cast(tf.cast(input_shape[0], tf.float32) * random_scale_factor,
        tf.int32)
    scaled_width = tf.cast(tf.cast(input_shape[1], tf.float32) * random_scale_factor,
        tf.int32)
```

```

scaled_shape = tf.squeeze(tf.stack([scaled_height, scaled_width]))
image = tf.image.resize(image, scaled_shape)
image = tf.image.random_crop(image, input_shape, seed=seed)
else:
    image = tf.image.resize(image, input_shape[:2])
label = features['label']
features = (image, label)
return features

train_data = train_data.map(normalizer)
val_data = val_data.map(normalizer)

print(train_data)
print(val_data)
class_names = train_info.features["label"].names
print(class_names)

plt.figure(figsize=(10, 10))
for i, (image, label) in enumerate(train_data.take(24)):
    # image = image.numpy().reshape([28,28,3])
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(image, cmap=plt.cm.binary)
    plt.xlabel(class_names[label])
plt.show()

train_data = train_data.batch(BATCH_SIZE)
val_data = val_data.batch(BATCH_SIZE)
train_data = train_data.prefetch(1)
val_data = val_data.prefetch(1)
# train_data = train_data.cache()

print(val_data)
print(train_data)

Inception_url = "https://tfhub.dev/google/tf2-preview/inception_v3/feature_vector/2"
inception_v3 = hub.KerasLayer(
    Inception_url, trainable=False,
    input_shape=[299, 299, 3],
    output_shape=[2048],
    dtype=tf.float32
)

LeNet = tf.keras.Sequential([
    inception_v3,
    tf.keras.layers.Dense(10, activation='softmax', name='logits_pred')
], name="LeNet")

print(LeNet.summary())

```

```

model_dir = './models/LeNet'
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
callbacks = [
    # Callback to interrupt the training if the validation loss/metrics stops improving for s
    tf.keras.callbacks.EarlyStopping(patience=8, monitor='val_acc',
                                    restore_best_weights=True),
    # Callback to log the graph, losses and metrics into TensorBoard:

    tf.keras.callbacks.TensorBoard(model_dir, histogram_freq=1, write_graph=True)

    # Callback to simply log metrics at the end of each epoch (saving space compared to verbose
]
LeNet.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=[tf.keras.metrics.SparseCategoricalAccuracy(name='acc'),
                      tf.keras.metrics.SparseTopKCategoricalAccuracy(k=5, name='top5_acc')])
history = LeNet.fit(train_data, epochs=50, steps_per_epoch=train_steps_per_epoch, validation_data=

```



Downloading and preparing dataset **cifar10/3.0.2** (download: 162.17 MiB, generated: 132.4

DI Completed...: 100% 1/1 [00:04<00:00, 4.55s/ url]

DI Size...: 100% 162/162 [00:04<00:00, 35.82 MiB/s]

Extraction completed...: 100% 1/1 [00:04<00:00, 4.47s/ file]

Shuffling and writing examples to /root/tensorflow_datasets/cifar10/3.0.2.incompleteTZQ
87% 43270/50000 [00:00<00:00, 63820.75 examples/s]

Shuffling and writing examples to /root/tensorflow_datasets/cifar10/3.0.2.incompleteTZQ
0% 0/10000 [00:00<?, ? examples/s]


Dataset cifar10 downloaded and prepared to /root/tensorflow_datasets/cifar10/3.0.2. Sub
<PrefetchDataset shapes: {id: (), image: (32, 32, 3), label: ()}, types: {id: tf.string
40000

5000

<MapDataset shapes: ((299, 299, 3), ()), types: (tf.float32, tf.int64)>


<MapDataset shapes: ((299, 299, 3), ()), types: (tf.float32, tf.int64)>

['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truc





frog



truck



airplane



deer

```
<PrefetchDataset shapes: ((None, 299, 299, 3), (None,)), types: (tf.float32, tf.int64)>
<PrefetchDataset shapes: ((None, 299, 299, 3), (None,)), types: (tf.float32, tf.int64)>
Model: "LeNet"
```

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 2048)	21802784
logits_pred (Dense)	(None, 10)	20490
Total params: 21,823,274		
Trainable params: 20,490		
Non-trainable params: 21,802,784		

None

```
LeNet.save_weights("lenet_weights", overwrite=True)
```

```
LeNet.save('LeNet_with_augmentation')
```

```
INFO:tensorflow:Assets written to: LeNet_with_augmentation/assets
```

```
INFO:tensorflow:Assets written to: LeNet_with_augmentation/assets
```

```
epoch 4/50
```

```
test_data = tfds.load('cifar10', split='test')
```

```
1250/1250 [=====] - 1/0s 136ms/step - loss: 0.5770 - acc: 0.80
```

```
LeNet = tf.keras.models.load_model('LeNet_with_augmentation')
```

```
epoch 1/50
```

```
num_tests=0
```

```
for _ in test_data:
```

```
    num_tests+=1
```

```
print(num_tests)
```

```
10000
```

```
1250/1250 [-----] - 1/0s 136ms/step - loss: 0.3124 - acc: 0.80
```

```
test_data = test_data.map(normalizer)
```

```
test_data = test_data.batch(BATCH_SIZE)
```

```
test_data = test_data.prefetch(1)
```

```
Epoch 14/50
```

```
LeNet.evaluate(test_data, batch_size=BATCH_SIZE, verbose=1)
```

```
313/313 [=====] - 40s 125ms/step - loss: 0.4954 - acc: 0.8323
[0.4953974485397339, 0.8323000073432922, 0.9937000274658203]
```

```
1250/1250 [=====] - 1/0s 136ms/step - loss: 0.2782 - acc: 0.90
```

```
%load_ext tensorboard
```

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

Mounted at /content/drive

```
ls
```

```
checkpoint                lenet_weights.index      models/  
lenet_weights.data-00000-of-00001  LeNet_with_augmentation/ sample_data/
```

```
%tensorboard --logdir models/LeNet/
```